Can user facing and background functions coexist in serverless?

Anonymous Author(s)

1 Introduction

A world where cloud compute is run in the format of serverless functions is attractive to developers and providers: developers pay only for what they use, while having access to many resources when needed; and cloud providers have control over scheduling and can use that to drive up utilization, rather than needing to hold idle resources available for clients who reserved them.

There remain roadblocks, however, that make serverless today infeasible for workloads that are a good fit. A central example to this paper is that of web applications. Their inconsistent and bursty load patterns make them a great fit for serverless, but they are rarely run completely in serverless offerings [15, 33, 34]. One of the reasons is serverless function invocations' variable end to end latencies: in a small benchmark (described in Section 2) we found that total execution time latencies for a simple hello world function that sleeps for 20 ms ranged from 20 to 400ms in AWS. This variability is a problem because it has been shown that small response time differences can matter a lot in interactive applications [11, 17], so a maximum acceptable latency for a user-facing page might be more like 100ms [23].

A well-known cause of these variable latencies is cold starts. This paper takes the position that systems research is well underway to reaching low single digit ms cold start times, with current state-of-the-art research systems pushing into single digit ms territory [29, 32]. Which begs the question: if cold start is fast enought that latency sensitive applications like web applications can have a cold start on the critical path, will serverless then be ready to support these sorts of workloads?

This paper argues that there still remains a challenge to running such a latency sensitive workload on serverless: queueing and delay within the system. We call this the *big noisy neighbor problem*, or BNNP for short: when load is higher than the resources can process, the scheduler has to queue or delay some functions. This means that the runtime of functions during high load is determined by how much load other functions have, and how many resources they use (hency big noisy neighbors).

This paper proposes XX, a serverless scheduler that addresses the BNNP by ensuring that latency sensitive functions don't end up behind background ones, waiting to be placed on machines or to get access to resources. Designing XX faces multiple challenges.

One challenge is that XX needs to be able to support a multi-tenant environment. A core facet of the BNNP is that it arises in a multi-tenant setting: part of the reason it is a problem is because I can't control other peoples' load, and there is no fair reason that it should impact my functions performance.[hmng: awk] The BNNP is also more challenging in the multi-tenant setting: it is clear that user-facing functions are more important to a web application than image processing jobs, but there is no way of mapping that prioritization to other developers' jobs, in order to be able to directly compare which job should run when.

In order to achieve this global comparability, XX uses *price classes*: the amount that it costs to run functions per unit time. This connection to money allows the price class to have meaning on an absolute scale. It also incentivizes usage of the lower price classes for functions that are less latency sensitive, since they can be run cheaply.

Another challenge is that of placing functions quickly enough. Knowing where the free and idle resources are, or finding out quickly, is challenging in a setting where both the number of new functions invocations and the amount of resources are large.

A third key challenge in designing XX is that of managing memory. For compute resources, cores can be timshared or processes preempted, but the buck stops once a machine is out of memory. Current systems address this challenge by requiring developers to express a maximal amount of memory they will use, and charging based on that. However, memory usage is at best difficult to know in advance and at worst has a large variance so is impossible to say in advance. And more importantly, is not correlated with what developers actually care about, which is function latency. Instead, XX charges developers based on the amount of memory actually used, and requires no bound to be set. XX thus faces the challenging proposition of blindly placing functions not knowing how much memory they will use, but still needing memory utilization to be high.

2 Motivation

This paper is motivated by the benefits of serverless for workloads like web applications, and shows that the BNNP will need to be solved in order to achieve it.

Figure 1: distribution of end to end duration times. The y axis is log scale

2.1 Web applications are a good fit for serverless

Web applications' traffic patterns make them a great candidate for running as serverless functions: their load is event-based, bursty and unpredictable, and a function's resource requirements can vary greatly depending on which user invoked it.

A back of the envelope calculation shows that for web applications with small load, lambda functions today are cheaper: with 50K requests per day, a memory footprint of < 128 MB per function and 200ms of execution, running that on AWS lambda adds up to $1.58 per month. On the other hand, the cheapest EC2 instance costs just over $3 per month. Of course, as the number of requests goes up, the price for lambdas scales linearly, whereas running an EC2 instance on full load becomes comparably cheap. Extensive simulations show a more nuanced picture of the tradeoff points for different workloads [10, 31].

Serverless also may outperform reservation systems for workloads that are very bursty: starting a new lambda execution environment is much faster than starting a new container or EC2 instance, which can take multiple minutes [1].

2.2 The BNNP

Any system that runs with high average utilization must experience moments where there is more load than the resources can handle. As we know from recent traces [35], although load may look stable at a time increment of hours or even minutes, going down to the second level shows the load to have high variance. If providers want to have good average utilization, then in the moments of load spike the load will be more than the resources can handle.

In a multi tenant setting, what this means is that one developers exprienced latencies will depend on other people's load. This is not ideal, expecially if one developer's user facing function is queued while another developers background task is running. This situation is the essence of the BNNP.

We can see evidence of the BNNP occurring in lambda invocations today, within the cold start invocations. We run an experiment with a simple lambda function that sleeps for 20ms and then returns. We use AWS Xray [2] to measure its

latency, with incovations spaced randomly between 0 and 10 minutes. The results are in Figure 1. The spike on the left side of the graph is the execution times from invocations that used warm start. The durations remain stable, because AWS is able to simply route the new request to the machine with the existing container. We verify this what is happening by changing the function to include reading and then writing to an environment variable, and find that for invocations with warm start when we read the variable it was already set by a previous invocation.

The right grouping in the graph is those invocations that hit cold starts, whose overall latencies vary between ~200 and ~400ms. This variance indicates that there is something more going on than just waiting for a container to start.

Although AWS' scheduling mechanism is proprietary, we can look to open source alternatives to see what common best practices exist today. Schedulers have two different options when load exceeds compute capacity: queue the excess load, or place it on machines and let them deal with being temporarily overloaded. Different schedulers have different approaches.

In OpenWhisk [26], the load balancer will choose which machine to run the function on, and then place the invocation, adressed to that machine, into a Kafka queue that the machine subscribes to and can pull the invocation from when it is ready [3]. This means that in the case of high load, a latency sensitive function might sit in the Kafka queue while someone else's background function completes: the BNNP.

Knative [20] similarly queues the excess invocations, although it does so via the load balancer, which is also in charge of autoscaling [21]: if the existing pods are fully loaded (with a small, bounded-size queue in front of them), requests are queued separately while the autoscaler starts up more invocations. Again, latency sensitive functions are potentially waiting for someone else's background function completes: the BNNP.

Hermod [19] is a recent research serverless scheduler, and shows in a simulation that late binding (as openwhisk and knative do) performs worse than early binding. Under high load, Hermod places the excess functions on machines anyway (used a least loaded policy) and does Processor Sharing scheduling among them. This ensures that no one function has a high delay, but rather that all are equally slowed down. This still leads to the BNNP though, just that now rather than experiencing queueing the latency senstive function is experiencing delay. The paper also does not address what happens when the machines are out of memory.

Because none of these schedulers have information about the functions they are queueing, it is impossible for them to know which to prioritize. The way all of the above schedulers avoid the BNNP is by doing different forms of accounting

concurrency: concurrency can be reserved or provisioned for specific functions, and limited for others. This is necessary to ensure that a burst in background tasks doesn't starve the latency sensitive functions. However, what happens when the datacenter is out of resources and so the concurrency limit has not yet been reached but the resources are unavailable is not clear. Reserving and provisioning and limiting are also conceptually in tension with the goal of serverless, which is to be on-demand and flexible.

3 Approach & Design

This paper's approach addresses the BNNP by using price classes as a metric to tell XX what to prioritize and what not. We will show that having price classes allows XX to stabilize the runtimes for high price class functions.

XX uses price classes to supplant the current interface, which requires developers to choose an amount of memory per function, which is tied to a cpu fraction (e.g., 0.2 vCPUs). In XX, price classes are the only thing that developers need to give; XX bills memory separately and by use. The price for memory is the same across all price classes. Removing memory from the interface serves the purpose of extending the serverless on-demand structure to include memory.

Price classes are a metric that has a number of benefits over resource usage estimations. One is that developers are more likely to have a good sense of what price class a function should have ahead of time, because they know in what context the function will be used and how important it is that the function run quickly. Price classes also remain the same across different invocations, whereas the resources needed can be heavily skewed in a web application environment, where popularity distributions are often very uneven [19, 28]. And finally, price classes more directly align the interests of the developer with those of the provider, by communicating on the level of what the provider and developer actually care about: money, and latency (as achieved by price classes in the system).

However, having price classes also means that there are no absolute guarantees about what developers are receiving when they put a price on a function. In order to mitigate that and avoid the developer-side uncertainty of bidding wars, XX exposes a fixed set of price classes. This fixed price list is similar to how AWS has different EC2 instance types, that are directly mapped to prices. Rather than being a guarantee, the price class is instead a metric to express priority to XX, which XX uses to enforce a favoring of high price class functions.

Having price classes also allows the provider to provision their datacenters in terms of the amount of hardware they buy: by looking at the historical overall amount of high price class load, they know a minimum of how much hardware they need to buy to be able to comfortably fit that load.
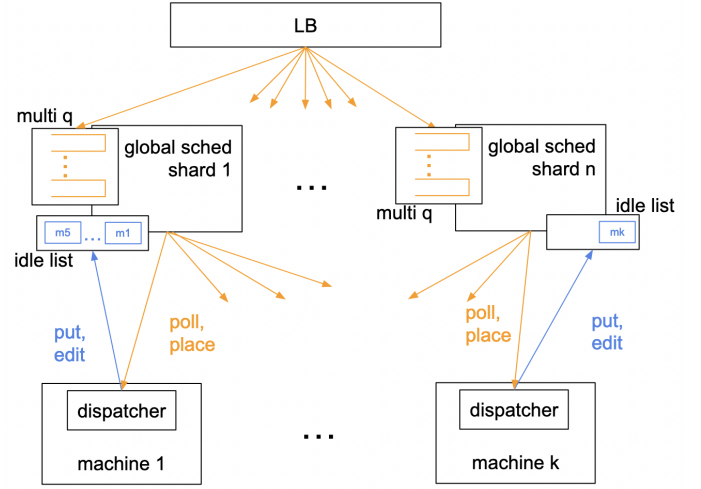

Figure 2: global scheduler shards queue and place functions (in orange), on each machine a dispatcher thread keeps track of memory utilization and if it's low writes itself to an idle list (in blue)

At the same time, price classes give XX the information it needs to decide what to schedule when. How exactly it does this is what we explore in XX's design.

3.1 Interface

Developers using XX write function handlers and define triggers just like they would for any existing serverless offering. In addition, each place where they trigger the function, they assign that invocation to a price class. For instance, a simple web application might consist of a home page view that is assigned a higher price class and costs $2\mu\dot{c}$ per cpu second, a user profile page view which is assigned a middle-high price class and cost $1.5\mu\dot{c}$ per cpu second, and finally an image processing function that can be set to a low price class which costs only $0.5\mu\dot{c}$ per cpu second.

Price classes are inherited across call chains: if a high price class function calls a low price class function, that invocation with run with high price class. This inheritance is important in order to avoid priority inversion.

To avoid unexpected costs in the case of for example a DOS attack or a bug, developers also express a monthly budget that they are willing to pay. XX uses this budget as a guideline and throttles invocations or decreases quality of service in the case that usage is not within reason given the expected budget, though it does not guarantee that the budget will not be exceeded by small amounts.

3.2 XX Design

XX has as its goal to enforce the price classes attached to functions, which means it needs to prefer higher price class

functions over lower ones, and preempt the latter when necessary.

As shown in Figure 2, XX sits behind a load balancer, and consists of: a *distributed global scheduler*, which places new function invocations and has attached an *idle list*, a *dispatcher*, which runs on each machine and communicates with the global scheduler shards, and a *machine scheduler*, which enforces price classes on the machines.

Machine Scheduler. The machine scheduler is a pre-emptive priority scheduler: it preempts lower price class functions to run higher price class ones. Being unfair and starving low price class functions is desirable in XX, since image processing functions should not interrupt a page view request processing, but vice versa is expected. Within price classes the machine scheduler is first come first served. This design matches Linux' 'SCHED FIFO' scheduling [4].

Idle list. Each global scheduler shard has an idle list, which holds machines that have a significant amount of memory available. In the shards idle list, each machine's entry is associated with the amount of memory available as well as the current amount of functions on the machine. The idle list exists because datacenters are large: polling a small number of machines has been shown to be very powerful, but cannot find something that is a very rare occurrence [24]. Having an idle list allows the machines that have actually idle resources, which are expected to be rare in a high-utilization setting, to make themselves visible to the global scheduler. The idle list also allows the global scheduler to place high price class functions quickly, without incurring the latency overheads of doing the polling to find available resources. This design is inspired by join idle queue [24], but defines idleness via memory availability rather than empty queues.

Dispatcher. The dispatcher is in charge of adding itself to a shard's idle list when memory utilization is low. The dispatcher chooses which list to add itself to using power-of-k-choices: it looks at k shards' idle lists and chooses the one with the least other machines in it. If the machine is already on an idle list on shard i, but the amount of available memory has changed significantly (either by functions finishing and memory being freed or by memory utilization increasing because of new functions or memory antagonists), the dispatcher will update shard i's idle list. These interactions from the dispatcher to free lists are represented by the blue arrows in Figure 2.

The dispatcher is also in charge of managing the machine's memory. When memory pressure occurs, the dispatcher uses *price class-based swapping* to move low price class functions off the machine's memory. Having priority scheduling creates an opportunity: because the dispatcher knows that the lowest price class functions will not be run until the high price class functions have all finished, it can swap its memory out knowing it will not be needed soon. The dispatcher swaps the low price class function back in when the memory pressure is gone and the function will be run.

Bounding the amount of swap space required without bounding the amount of memory that functions can use is not possible. The goal of the dispatcher is to swap when possible, and in the case that that is not enough it can resort to killing. Providers can estimate the amount of swap space required by looking at memory utilization, and since swap space is cheap [5] can provision it so that killing is very rare.

Global Scheduler Shards. Global scheduler shards store the functions waiting to be placed in a multi queue, with one queue per price class. Shards choose what function to place next by looking at each function at the head of each queue, and comparing the ratio of price class to amount of time spent in the queue. This procedure ensures that high price class functions don't have to wait as long as low price class functions to be chosen next, but low price class functions will get placed if they have waited for a while.

When placing the chosen function, the shard will first look in its idle list. If the list is not empty, it will choose the machine with the smallest queue length.

If there are no machines in the idle list, the shard switches over to power-of-k-choices: it polls k machines, getting the amount of functions running from each. The shard then places the new function on the machine with the smallest number of currently running functions.

4 Preliminary Results

In order to explore evidence for the case for XX, we ask the following questions:

(1) How does function latency in XX compare to schedulers without priorities?
(2) Does XX's plan for managing memory work?

To explore these questions, we build a simulator in go[6], which simulates different scheduling approaches. Using a simulator allows us to extend the experiments to include many more machines than would otherwise be available to us.

4.1 Experimental Setup

In each version of the simulator, functions arrive in an open loop at a given rate. The simulator attaches three main characteristics to each function it generates: runtime, price class, and memory usage. *Function runtime* is chosen by sampling from randomly generated long tailed (in this case pareto) distribution: the relative length of the tail (α value) remains constant, and the minimum value (x_m) is chosen from a normal distribution. This process for sampling reflects the fact that different functions have different expected runtimes (chosen from a normal distribution), and that actual invocation runtimes follow long tailed distributions (so each pareto distribution that we sample represents the expected runtime
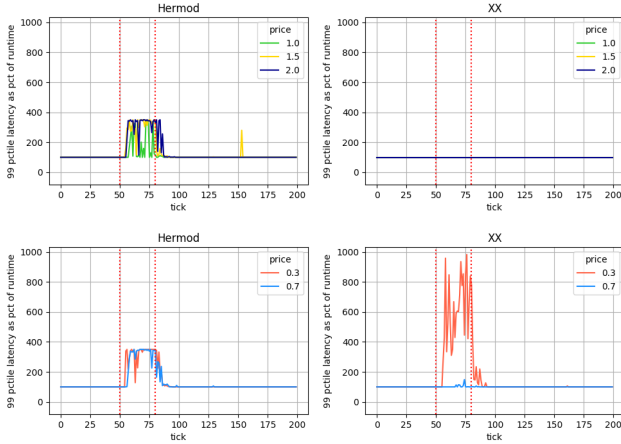
Figure 3: tail latency distribution and compute utilization for Hermod and XX, for high (top) and low (bottom) price class functions. At tick 50 (left right line) the load was increased, and at tick 80 (right red line) it was decreased again

distribution of a given function). *Function price class* is chosen randomly, but weighted: the simulator uses a bimodal weighting across priorities. The simulator has n different price class values, each assigned to a fictitious price. Because functions are randomly assigned a price class, runtime and price class are not correlated. *Function memory usage* is chosen randomly between 100MB and 10GB. Over their lifetime, functions increase their memory usage from an initial amount (always 100MB) to their total usage.

When comparing two different simulated schedulers, they each are given an identical workload and then each simulate running that workload.

The simulator makes some simplifying assumptions:

(1) functions are compute bound, and do not block for i/o
(2) communication latencies are not simulated
(3) the amount of time it takes to swap memory is not simulated

We simulate running 100 machines with 8 cores each, 4 scheduler shards, and a k-choices value of 3 when applicable.

4.2 How do function latencies compare?

The goal of XX is to reduce the variance of latencies for high price class functions. This experiment shows that XX is able to do this even under varying load, and compares XX's performance to a scheduler that does not take into account any information about which functions are latency sensitive.

We simulate Hermod[19], a state-of-the-art research scheduler built specifically for serverless. Hermod's design is the result of a from-first-principles analysis of different scheduling paradigms. In accordance with the paper's findings, we simulate least-loaded load balancing over machines found using power-of-k-choices, combined with early binding and

Processor Sharing machine-level scheduling. Hermod does not use priorities in its design, and as such the simulator ignores functions' price class when simulating Hermod's design.

Because Hermod does not deal with memory pressure, and to avoid an unfair comaprison with XX's swapping, we set the memory to be absurdly high for all XX as well as Hermod in this experiment. We also turn off the use of the idle list in XX, so as to be on par with Hermod in placing load, and revert solely to k-choices.

We begin with a medium load setting, and then increase the load to more than a steady setting cand handle for a window of time. A strong result for XX would show that it is able to maintain low latency for high price class functions, even under the high load. Figure 3 shows the results. We can see that XX is indeed able to maintain low latencies for the high price class functions. Hermod spreads the performance degradation across all the different functions, and as a result all of their latencies go up equally.

4.3 Does XX plan for memory management work?

To answer this question, we look at how XX distributes load, and whether the amount that XX needs to swap memory is realistic. We now run XX in a setting of limited memory (32GB of RAM per machine), and track the memory utilization of different machines, as well as how much and what machines need to swap. A good result would show: a small spread of memory utilization, that machines only start swapping once memory utilization is high, and that the amount of swapping being done is equally spread across machines. Figure 4 shows the results. We can see XX swaps only lower price class functions' memory, and that the amount of memory swapped is faily evenly distributed between all the machines. We can also conclude that with a 500GB SSD, a provider would comfortably be able to avoid killing while running the datacenter at an average memory utilization of ∼90%, at the cost of ∼ \$30 per machine for swap space [5].

5 Related Work

Many other projects have explored how to do better scheduling for data centers.

Systems like Sparrow[27], Hermod[19], or Kairos[14] improve performance of scheduling in the distributed setting by trying out and using different scheduling policies. Unlike XX, they treat all functions equally.

Like XX, many projects tailor their approach to serverless. Some systems generate information about functions themselves to help placement decisions; for instance ALPS[16], which observes and learns the behaviors of existing functions and then makes scheduling decisions based on those; or Morpheus[18], which learns SLOs from historical runs, and then translates these to recurring reservations. XX instead
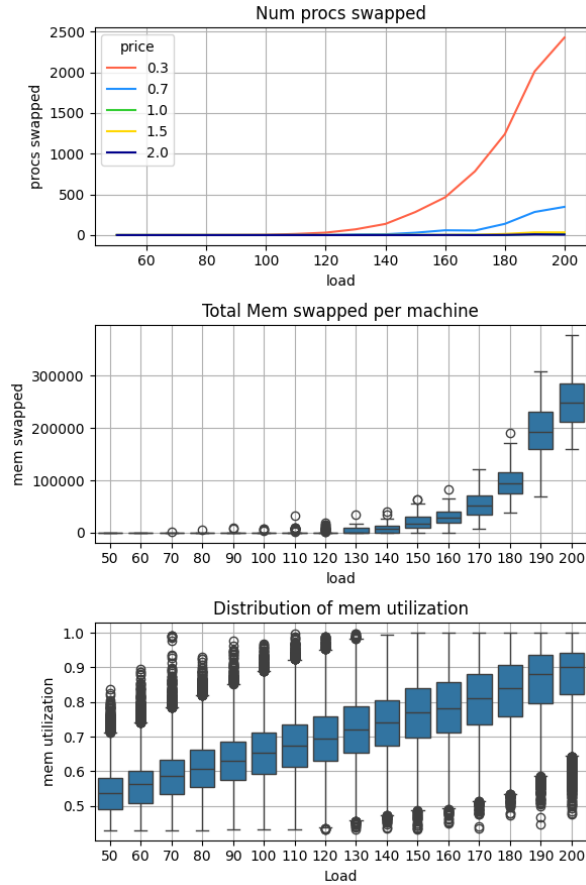
Figure 4: XX's swapping behavior. The amount of memory is in MB

gets the price classes directly from the developers as part of its interface.

Other papers have taken the same approach as XX of getting information to help scheduling from the developers. Sequoia[30], for instance, creates a metric of QOS for serverless functions. Unlike XX however, Sequoia does not implement a new scheduler, but is itself a serverless function that manages the invocation sequence of developer's function chains by interposing on the triggers and choosing what to invoke when. Therefore it also, unlike XX, does not support multi-tenancy.

Allocation Priority Policies (APP)[13] provides a declarative language to express policies, then builds a scheduler around that. The APP language is built around allowing developers to specify custom load balancing decisions, and the scheduler uses the developers' specification to define a mapping of function invocations to workers. XX, on the other hand, does not ask developers to set the load balancing policy, but rather has developers give XX the information it needs to do the load balancing itself.

AWS offers two different ways for developers to influence their functions' scaling: provisioned and reserved concurrency[7]. Provisioned concurrency specifies a number of instances to keep warm for a given function, and reserved concurrency ensures that a fixed amount of the possible concurrency reserved for it. This interface is bad for serverless workloads for the same reasons that reservation-based systems are: it requires developers to estimate their future needs and pay up front, and providers to keep those potentially idle resources available.

Serverless orchestration systems like Dirigent [12] are orthogonal to XX: thier approaches can be combined to further reduce the latency overheads that functions face.

On the serverful side of scheduling, priorities are generally expressed via a latency critical/best effort binary, where latency critical processes have an attached amount of resource reservations, and best effort processes don't [22]. Serverless schedulers that sit on top of Kubernetes, such as OpenFaaS [25] or Knative [20], use autoscaling to keep up with load, and use the same interface as kubernetes for developers to express resource requirements [8, 9]. This works well for long running servers with steady amounts of load, since predictable load will allow developers to make good approximations of the resources they will need. The serverless setting XX works within is different because both the number of invocations and the resource usage of each invocation is expected to vary.

6 Conclusion

Serverless was and is a great option for developers whose load varies and providers who don't want to keep resources idle for processes that reserved them. However, the reality of serverless today is that functions experience a variance in latencies that is not tolerable for latency sensitive workloads.

We propose a new scheduler, XX, that introduces *price classes*. Developers assign each function they want to run to a price class, which encodes a priority that XX then enforces at invocation, both in placing the function and on the machine level by running priority scheduling. We show that XX is able to enforce priorities and keep high price class functions latencies stable even under high load.

References

[1]   URL: https://docs.aws.amazon.com/autoscaling/ec2/userguide/ec2-auto-scaling-default-instance-warmup.html.

[2]   URL: https://aws.amazon.com/xray/.

[3]   URL: https://github.com/apache/openwhisk/blob/master/docs/about.md.

[4]   URL: https://man7.org/linux/man-pages/man7/sched.7.html.

[5]   URL: https://www.bestbuy.com/site/pny-cs900-500gb-internal-ssd-sata/6385542.p?skuId=6385542.

[6]   URL: https://go.dev/.

[7]   URL: https://docs.aws.amazon.com/lambda/latest/dg/lambda-concurrency.html.

[8]   URL: https://knative.dev/docs/serving/services/configure-requests-limits-services/.

[9]   URL: https://docs.openfaas.com/reference/yaml/.

[10]   Álvaro Alda Rodríguez et al. *Economics of 'Serverless'*. URL: https://www.bbva.com/en/innovation/economics-of-serverless/.

[11]   Jake Brutlag. *Speed Matters*. URL: https://research.google/blog/speed-matters/.

[12]   Lazar Cvetković, François Costa, Mihajlo Djokic, Michal Friedman, and Ana Klimovic. "Dirigent: Lightweight Serverless Orchestration". In: *Proceedings of the ACM SIGOPS 30th Symposium on Operating Systems Principles*. SOSP '24. Austin, TX, USA: Association for Computing Machinery, 2024, pages 369–384.

[13]   Giuseppe De Palma, Saverio Giallorenzo, Jacopo Mauro, and Gianluigi Zavattaro. "Allocation Priority Policies for Serverless Function-Execution Scheduling Optimisation". In: *Service-Oriented Computing: 18th International Conference, ICSOC 2020, Dubai, United Arab Emirates, December 14–17, 2020, Proceedings*. Dubai, United Arab Emirates: Springer-Verlag, 2020, pages 416–430.

[14]   Pamela Delgado, Diego Didona, Florin Dinu, and Willy Zwaenepoel. "Kairos: Preemptive Data Center Scheduling Without Runtime Estimates". In: *Proceedings of the ACM Symposium on Cloud Computing*. SoCC '18. Carlsbad, CA, USA: Association for Computing Machinery, 2018, pages 135–148.

[15]   Jesse Duffield. *My notes from deciding against AWS Lambda*. URL: https://jesseduffield.com/Notes-On-Lambda/.

[16]   Yuqi Fu, Ruizhe Shi, Haoliang Wang, Songqing Chen, and Yue Cheng. "ALPS: An Adaptive Learning, Priority OS Scheduler for Serverless Functions". In: *2024 USENIX Annual Technical Conference (USENIX ATC 24)*. Santa Clara, CA: USENIX Association, July 2024, pages 19–36.

[17]   Gigaspaces. *Amazon Found Every 100ms of Latency Cost them 1 Percent in Sales*. URL: https://www.gigaspaces.com/blog/amazon-found-every-100ms-of-latency-cost-them-1-in-sales.

[18]   Sangeetha Abdu Jyothi, Carlo Curino, Ishai Menache, Shravan Matthur Narayanamurthy, Alexey Tumanov, Jonathan Yaniv, Ruslan Mavlyutov, Inigo Goiri, Subru Krishnan, Janardhan Kulkarni, and Sriram Rao. "Morpheus: Towards Automated SLOs for Enterprise Clusters". In: *12th USENIX Symposium on Operating Systems Design and Implementation (OSDI 16)*. Savannah, GA: USENIX Association, Nov. 2016, pages 117–134.

[19]   Kostis Kaffes, Neeraja J. Yadwadkar, and Christos Kozyrakis. "Hermod: principled and practical scheduling for serverless functions". In: *Proceedings of the 13th Symposium on Cloud Computing*. SoCC '22. San Francisco, California: Association for Computing Machinery, 2022, pages 289–305.

[20]   *Knative*. URL: https://knative.dev/.

[21]   Stavros Kontopoulos. *Demystifying Activator on the data path*. URL: https://knative.dev/blog/articles/demystifying-activator-on-path/.

[22]   Kubernetes. *Configure Quality of Service for Pods*. URL: https://kubernetes.io/docs/tasks/configure-pod-container/quality-service-pod/.

[23]   Greg Linden. *Marissa Mayer at Web 2.0*. URL: http://glinden.blogspot.com/2006/11/marissa-mayer-at-web-20.html.

[24]   Yi Lu, Qiaomin Xie, Gabriel Kliot, Alan Geller, James R. Larus, and Albert Greenberg. "Join-Idle-Queue: A novel load balancing algorithm for dynamically scalable web services". In: *Performance Evaluation* 68.11 (2011). Special Issue: Performance 2011, pages 1056–1071.

[25]   *OpenFaaS*. URL: https://www.openfaas.com/.

[26]   *OpenWhisk*. URL: https://openwhisk.apache.org.

[27]   Kay Ousterhout, Patrick Wendell, Matei Zaharia, and Ion Stoica. "Sparrow: distributed, low latency scheduling". In: *Proceedings of the Twenty-Fourth ACM Symposium on Operating Systems Principles*. SOSP '13. Farminton, Pennsylvania: Association for Computing Machinery, 2013, pages 69–84.

[28]   Mohammad Shahrad, Rodrigo Fonseca, Inigo Goiri, Gohar Chaudhry, Paul Batum, Jason Cooke, Eduardo Laureano, Colby Tresness, Mark Russinovich, and Ricardo Bianchini. "Serverless in the Wild: Characterizing and Optimizing the Serverless Workload at a Large Cloud Provider". In: *2020 USENIX Annual Technical Conference (USENIX ATC 20)*. USENIX Association, July 2020, pages 205–218.

[29]   Ariel Szekely, Adam Belay, Robert Morris, and M. Frans Kaashoek. "Unifying serverless and microservice workloads with SigmaOS". In: *Proceedings of the ACM SIGOPS 30th Symposium on Operating Systems Principles*. SOSP '24. Austin, TX, USA: Association for Computing Machinery, 2024, pages 385–402.

[30]   Ali Tariq, Austin Pahl, Sharat Nimmagadda, Eric Rozner, and Siddharth Lanka. "Sequoia: enabling quality-of-service in serverless computing". In: *Proceedings of the 11th ACM Symposium on Cloud Computing*. SoCC '20. Virtual Event, USA: Association for Computing Machinery, 2020, pages 311–327.

[31]   Andy Warzon. *AWS Lambda Pricing in Context - A Comparison to EC2*. URL: https://www.trek10.com/blog/lambda-cost.

[32]   Xingda Wei, Fangming Lu, Tianxia Wang, Jinyu Gu, Yuhan Yang, Rong Chen, and Haibo Chen. "No Provisioned Concurrency: Fast RDMA-codesigned Remote Fork for Serverless Computing". In: *17th USENIX Symposium on Operating Systems Design and Implementation (OSDI 23)*. Boston, MA: USENIX Association, July 2023, pages 497–517.

[33]   *Why or why not use AWS Lambda instead of a web framework for your REST APIs?* URL: https://www.reddit.com/r/Python/comments/1092py3/why_or_why_not_use_aws_lambda_instead_of_a_web/.

[34]   *Without saying "it's scalable", please convince me that a serverless architecture is worth it.* URL: https://www.reddit.com/r/aws/comments/yxyyk3/without_saying_its_scalable_please_convince_me/.

[35]   Bartek Wydrowski, Robert Kleinberg, Stephen M. Rumble, and Aaron Archer. "Load is not what you should balance:

Introducing Prequal". In: *21st USENIX Symposium on Networked Systems Design and Implementation (NSDI 24)*. Santa Clara, CA: USENIX Association, Apr. 2024, pages 1285–1299.